


REVENUE Pre-revenue no ARR/KPIs	VALUATION ~\$500M post-money	TOTAL RAISED \$75M
EMPLOYEES ~6 as of the April 2026 Series A		

BUSINESS OVERVIEW

- Standard Intelligence (legally Standard Intelligence PBC; si.inc) is a San Francisco AI research lab building general-purpose foundation models for 'computer use', agentic AI that operates software directly through its graphical interface the way a human does.
- Its flagship model, FDM-1, is trained on an ~11-million-hour corpus of raw screen-recording video; the team argues video pretraining is the only way to scale the 'action data' general agents need.
- The lab is famously tiny, roughly six people, and in April 2026 raised a \$75M Series A led by Sequoia Capital and Spark Capital at a reported ~\$500M post-money valuation, with angels Andrej Karpathy, Stanley Druckenmiller, and Tesla Optimus lead Milan Kovac.
- It first earned research credibility by open-sourcing 'hertz-dev' (Nov 2024), an 8.5B-parameter full-duplex conversational audio model.
- The company is pre-revenue: the bet is on a contrarian architecture and a long-horizon path to aligned general intelligence, not near-term monetization.

How the business works, left to right: who buys, what the platform provides, and the payoff.



SOURCES: SILICONANGLE — STANDARD INTELLIGENCE RAISES \$75M · SEQUOIA CAPITAL — TRAINING GENERAL INTELLIGENCE IN PIXEL SPACE ·

THE INFORMATION — STANDARD INTELLIGENCE RIDES NEOLAB FERVOR

EXECUTIVE LEADERSHIP

CO-FOUNDER (CEO)

Galen Mead

Co-founded Standard Intelligence at age ~21; met co-founder Devansh Pandey as a teenager through the Atlas Fellowship (2022). Left undergraduate studies to build the company. Credited contributor on hertz-dev and FDM-1.

CO-FOUNDER

Devansh Pandey

Co-founded Standard Intelligence at age ~20; met Galen Mead via the Atlas Fellowship (2022); left his undergraduate program to pursue the company's video-pretraining approach to general agents. Credited contributor on FDM-1 and hertz-dev.

RESEARCHER / FOUNDING TEAM MEMBER

Neel Redkar

Listed as a researcher on FDM-1 and credited on hertz-dev; part of the ~6-person San Francisco team.

RESEARCHER / FOUNDING TEAM MEMBER

Yudhister Kumar

Listed as a researcher on the FDM-1 model among the ~6-person Standard Intelligence team in San Francisco.

PRODUCT & REVENUE MODEL

- FDM-1 is a 'computer action model' that learns to operate software from video rather than from annotated screenshots or text.
- The team trained an inverse dynamics model (a masked-diffusion network) on ~40,000 hours of labeled footage to automatically infer keystrokes and mouse movements from on-screen changes, then used it to auto-label the full ~11M-hour video corpus, moving computer use from a data-constrained to a compute-constrained regime.
- A custom video encoder is claimed to be ~50x more token-efficient than prior state of the art (and ~100x vs. OpenAI's), fitting roughly two hours of 30 FPS footage in a 1M-token context window.
- The forward dynamics model predicts the next action directly from video and action tokens, no chain-of-thought, BPE tokenization, or tool use.
- Demonstrated (research-stage, not productized) capabilities include extruding a CAD gear in Blender, autonomously finding a banking-app bug via state-space 'fuzzing,' and learning to physically drive a Toyota RAV4 around a San Francisco block after under an hour of fine-tuning.
- There is no commercial product or revenue model yet.

SOURCES: STANDARD INTELLIGENCE — FDM-1 TECHNICAL WRITE-UP ·
SILICONANGLE — STANDARD INTELLIGENCE RAISES \$75M ·
SEQUOIA CAPITAL — TRAINING GENERAL INTELLIGENCE IN PIXEL SPACE

RECENT NEWS

2026-04-30

Standard Intelligence raises \$75M Series A from Sequoia and Spark Capital to build computer-use modelsSOURCE: [SILICONANGLE](#)

Validates a ~6-person team at a reported ~\$500M post-money, signaling intense investor appetite for video-trained computer-use models versus screenshot/API-based agent approaches.

2026-04-30

Standard Intelligence Rides Neolab Fervor with Computer Use ModelSOURCE: [THE INFORMATION](#)

The Information frames Standard as part of a new wave of well-funded 'neolab' AI startups and reports FDM-1 being road-tested (a RAV4 with a dashboard laptop) in San Francisco, real-world ambition beyond benchmarks.

2026-04-30

A Brief Fundraising Update: \$75M with angels Karpathy, Druckenmiller and KovacSOURCE: [STANDARD INTELLIGENCE \(SI.INC BLOG\)](#)

Company-confirmed round naming Sequoia's Sonya Huang and Spark's Yasmin Razavi plus marquee angels, a strong syndicate-quality signal for a pre-revenue lab.

2026-02-23

Standard Intelligence unveils FDM-1, the first fully general computer action model trained on ~11M hours of videoSOURCE: [STANDARD INTELLIGENCE \(SI.INC BLOG\)](#)

FDM-1 is the company's flagship technical differentiator (an encoder ~50, 100x more token-efficient than prior SOTA) and the asset that anchored the subsequent Series A.

2026-02-23

FDM-1 autonomously drives a car through San Francisco after under one hour of fine-tuningSOURCE: [STANDARD INTELLIGENCE \(SI.INC BLOG\)](#)

Demonstrates strong transfer from the video-pretrained foundation model to a physical control task, the headline proof point that the architecture generalizes beyond on-screen software.

2024-11-03

Standard Intelligence open-sources hertz-dev, the first base model for full-duplex conversational audioSOURCE: [GITHUB \(STANDARD-INTELLIGENCE/HERTZ-DEV\)](#)

The team's earliest public credibility marker (an 8.5B Apache-2.0 audio model at ~120ms real-world latency on one RTX 4090) established the research pedigree that later attracted Sequoia and Spark.

FINANCIALS & KEY METRICS

Research & Traction Snapshot (pre-revenue) · As of 2026-06-04

LAST VALUATION ~\$500M post-money reported	CAPITAL RAISED \$75M Series A Apr 2026	TEAM SIZE ~6 people	LEAD INVESTORS Sequoia · Spark Capital
FLAGSHIP MODEL FDM-1 computer-use	TRAINING CORPUS ~11M hours of video	REVENUE Pre-revenue no ARR/KPIs	FOUNDED 2024 · San Francisco

- AI / Product:** The entire thesis is architectural and contrarian: rather than the language-plus-screenshots-plus-tool-calls scaffolding OpenAI and Anthropic use, Standard Intelligence trains a model to predict 'the next mouse movement, click, and keystroke from the pixels in front of it', explicitly 'the Tesla FSD approach applied to knowledge work.' The key claimed unlock is reframing computer use from a data-constrained to a compute-constrained problem via an inverse-dynamics model that auto-labels an ~11M-hour video corpus, paired with a video encoder ~50, 100x more token-efficient than prior SOTA. If the bet is right, generality emerges from scale and the model compounds across domains (CAD, finance, engineering, ML research) from one foundation; if it is wrong, video pretraining may never reach the reliability bar that brittle-but-shipping agent harnesses already clear.
- Demand:** There is no customer or revenue traction, the company is pre-product, so 'demand' is demonstrated capability plus investor demand, not bookings. FDM-1's marquee proof points are autonomously driving a Toyota RAV4 around an SF block after under an hour of fine-tuning, basic Blender CAD via continuous mouse movements, and autonomous bug-finding (fuzzing a banking app into a negative balance). The strongest market signal is the financing itself: a \$75M round at a reported ~\$500M valuation from Sequoia and Spark with Karpathy, Kovac, and Druckenmiller as angels, underwriting a six-person team almost entirely on research conviction and the prior credibility of the open-source hertz-dev model.

SOURCES: STANDARD INTELLIGENCE — FDM-1 TECHNICAL WRITE-UP ·
SEQUOIA CAPITAL — TRAINING GENERAL INTELLIGENCE IN PIXEL SPACE ·
SILICONANGLE — STANDARD INTELLIGENCE RAISES \$75M

KEY RISKS & DEBATES

- Pre-revenue, pre-product:** no commercial product, customers, or revenue exist. The ~\$500M reported valuation rests entirely on research conviction and demos. Any slip in reaching a reliable, shippable computer-use model raises material down-round risk.

- !

Competing against giants: OpenAI, Anthropic, and Google DeepMind are all shipping computer-use agents with vastly more compute, capital, distribution, and talent. A six-person, \$75M lab is structurally out-resourced if the incumbents' screenshot-plus-tools approach proves 'good enough.'
- !

Unproven core thesis: video pretraining in 'pixel space' is elegant but unvalidated at the reliability bar real automation demands. Impressive one-off demos (driving a car, a Blender gear) are not the same as the high-success-rate, low-variance execution enterprises require, the gap from demo to dependable agent is where most computer-use efforts stall.
- !

Compute and capital intensity: the strategy explicitly bets on scaling compute. \$75M is small relative to frontier training budgets; the lab will likely need a much larger growth round, and its burn/compute needs make it sensitive to GPU access, pricing, and the AI funding climate.
- !

Key-person and talent risk: value is concentrated in two young founders and a handful of researchers. Deep-pocketed labs aggressively poach exactly this profile; the loss of one or two people would be disproportionately damaging at this scale.
- !

Safety, liability, and regulatory exposure: autonomous agents that act on real software (and real cars) carry concrete safety and liability surface, an agent that 'fuzzes' a banking app to a negative balance illustrates both the capability and the hazard. Computer-use autonomy is likely to attract regulatory and enterprise-security scrutiny as it matures.
- !

Valuation vs. stage: ~\$500M post-money for a pre-revenue, ~6-person team prices in extraordinary execution. Comparable reputation-driven labs (Thinking Machines \$12B, SSI \$32B) show the category can support large marks, but also that future rounds will demand visible technical milestones to justify step-ups.

SOURCES: THE INFORMATION — STANDARD INTELLIGENCE RIDES NEOLAB FERVOR ·
SILICONANGLE — STANDARD INTELLIGENCE RAISES \$75M · STANDARD INTELLIGENCE — FDM-1 TECHNICAL WRITE-UP

COMPARABLE COMPANIES

COMPANY	CATEGORY	LAST VALUATION	LAST RAISE	NOTE
Standard Intelligence	Subject	~\$500M (Apr 2026)	\$75M Series A (Apr 2026)	Pre-revenue, ~6 people; video-pretrained computer-use model (FDM-1). Reported ~\$500M post-money (The Information).
Thinking Machines Lab	Frontier AI lab (Mira Murati)	\$12B (Jul 2025)	\$2B seed (Jul 2025)	Closest structural peer: pre-product, reputation-driven lab. Shows the mark a top-founder pre-product lab can command.

COMPANY	CATEGORY	LAST VALUATION	LAST RAISE	NOTE
Safe Superintelligence (SSI)	Frontier AI lab (Ilya Sutskever)	\$32B (Apr 2025)	\$2B Apr 2025	Another pre-product lab valued on founder reputation; backers include a16z, Greenoaks, Alphabet, Nvidia.
Cognition (Devin)	Agentic AI / coding agent	\$10.2B (Sep 2025)	\$400M Sep 2025	Closest 'agent that uses software' comp — but revenue-generating, unlike the subject's pre-product stage.
Mistral AI	Frontier AI lab (Europe)	~\$14B (Sep 2025)	€1.7B Series C (Sep 2025)	Europe's largest-ever venture round (ASML-led); a scaled, productized frontier-lab benchmark.
Anthropic	Frontier AI lab (Claude)	\$380B (Feb 2026)	\$30B Series G (Feb 2026)	Scaled-incumbent ceiling and a direct computer-use competitor; reportedly in talks at a >\$900B valuation.
OpenAI	Frontier AI lab (ChatGPT)	\$852B (Mar 2026)	\$122B round (Mar 2026)	Category top-end and a direct computer-use (Operator) competitor; largest private financing in history.

SOURCES: Private valuations and rounds from the cited reporting (TechCrunch, Bloomberg, CNBC, Sacra, PitchBook). These are last-round marks, not market prices, and the dates differ by company.

FUNDING HISTORY

SERIES A

\$75M

~\$500M (reported)

2026-04

Sequoia Capital, Spark Capital

Total raised: \$75M

NOTABLE INVESTORS

- Sequoia Capital (Sonya Huang)
- Spark Capital (Yasmin Razavi)
- Andrej Karpathy (angel/advisor)
- Stanley Druckenmiller (angel/advisor)
- Milan Kovac (angel/advisor, Tesla Optimus lead)

IPO READINESS

- Very early, pre-revenue research lab, ~6 employees, founded ~2024 by Galen Mead and Devansh Pandey.
 - An IPO is years away and not on the horizon; the company is at the foundational-research / pre-product stage.
 - The relevant near-term milestones are model reliability, a first commercial product, and a likely larger growth round, not public-market readiness.
-

SLIDE-READY BULLETS

- 1 Standard Intelligence is a ~6-person San Francisco AI lab building general-purpose 'computer use' foundation models; its FDM-1 model acts directly on screen pixels, predicting the next mouse move, click, and keystroke, rather than reasoning over screenshots and tool calls like OpenAI's Operator or Anthropic's Claude computer use.
 - 2 Raised a \$75M Series A in April 2026 led by Sequoia and Spark Capital at a reported ~\$500M post-money, with angels Andrej Karpathy, Stanley Druckenmiller, and Tesla Optimus lead Milan Kovac, an unusually high-pedigree syndicate for a pre-revenue, six-person team.
 - 3 The technical wedge is data scale: an inverse-dynamics model auto-labels an ~11M-hour video corpus (reframing computer use from data-constrained to compute-constrained), paired with a video encoder claimed ~50, 100x more token-efficient than prior SOTA, 'the Tesla FSD approach applied to knowledge work.'
 - 4 Proof points are demos, not products: FDM-1 drove a Toyota RAV4 around an SF block after <1hr of fine-tuning, did basic Blender CAD, and autonomously found a banking-app bug, impressive generalization, but pre-revenue with no commercial product or customers.
 - 5 Key debate: is a contrarian video-pretraining bet by a tiny lab the path to general computer agents, justifying a ~\$500M mark against pre-product peers like Thinking Machines (\$12B) and SSI (\$32B), or will incumbents with orders of magnitude more compute and distribution reach 'good enough' computer use first?
-

SMART DILIGENCE QUESTIONS

- Q1 What is the measured reliability of FDM-1 on standardized computer-use benchmarks (e.g., success rate and variance on OSWorld / WebArena-style tasks) versus OpenAI Operator and Anthropic's computer use, and how does success rate degrade on long-horizon, multi-step tasks where most agent harnesses fail? Demos show capability; what does the failure distribution look like?
 - Q2 What is the compute budget and roadmap the \$75M actually funds, how large a training run does FDM-1's next generation require, what GPU access is secured, and at what burn rate does the company need to raise again? Given frontier-scale compute costs, is \$75M enough to validate the thesis before the next round?
-

Q3

What is the path from research demos to a commercial product and durable moat, which first use case (CAD, software QA, finance ops) does the team productize, what data or distribution advantage compounds over time, and what prevents a better-capitalized incumbent from replicating the video-pretraining approach once it is published?

SLIDE KIT

Logo, stats, and copy-paste artifacts

Standard Intelligence

PRIVATE · ~\$500M post-money
(Series A, April 2026; reported by The Information — company did not confirm a figure)

si.inc

- Artificial Intelligence
- Foundation Models
- Computer-Use Agents

COPY-PASTE STATS

Pre-revenue (no ARR/KPIs)

Valued at ~\$500M post-money (Series A, April 2026; reported by The Information — company did not confirm a figure)

\$75M total raised

BRAND COLORS



#116600

BRAND TAGS

- Artificial Intelligence
- Foundation Models
- Computer-Use Agents

TOP 3 SLIDE BULLETS

- Standard Intelligence is a ~6-person San Francisco AI lab building general-purpose 'computer use' foundation models; its FDM-1 model acts directly on screen pixels, predicting the next mouse move, click, and keystroke, rather than reasoning over screenshots and tool calls like OpenAI's Operator or Anthropic's Claude computer use.
- Raised a \$75M Series A in April 2026 led by Sequoia and Spark Capital at a reported ~\$500M post-money, with angels Andrej Karpathy, Stanley Druckenmiller, and Tesla Optimus lead Milan Kovac, an unusually high-pedigree syndicate for a pre-revenue, six-person team.
- The technical wedge is data scale: an inverse-dynamics model auto-labels an ~11M-hour video corpus (reframing computer use from data-constrained to compute-constrained), paired with a video encoder claimed ~50, 100x more token-efficient than prior SOTA, 'the Tesla FSD approach applied to knowledge work.'

Sources: Company IR, earnings transcripts, Bloomberg, SEC filings, Sacra, Crunchbase, yfinance
DRAFT for internal use only. Human review required before distribution.

ATLAS